

Analysis of urban versus rural living

```
In [2]: #importing Libery
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import copy
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

# styling
pd.set_option('display.max_columns',150)
plt.style.use('bmh')
from IPython.display import display
%matplotlib inline
```

Data Collection

```
In [3]: #loading dataset in pandas dataset
data = pd.read_csv('responses.csv')
```

In [4]: *#check first five rows of the dataset*
`data.head()`

Out[4]:

	Music	Slow songs or fast songs	Dance	Folk	Country	Classical music	Musical	Pop	Rock	Metal or Hardrock	Punk	I
0	5.0	3.0	2.0	1.0	2.0	2.0	1.0	5.0	5.0	1.0	1.0	
1	4.0	4.0	2.0	1.0	1.0	1.0	2.0	3.0	5.0	4.0	4.0	
2	5.0	5.0	2.0	2.0	3.0	4.0	5.0	3.0	5.0	3.0	4.0	
3	5.0	3.0	2.0	1.0	1.0	1.0	1.0	2.0	2.0	1.0	4.0	
4	5.0	3.0	4.0	3.0	2.0	4.0	3.0	5.0	3.0	1.0	2.0	

In [5]: *#check last five rows of the dataset*
`data.tail()`

Out[5]:

	Music	Slow songs or fast songs	Dance	Folk	Country	Classical music	Musical	Pop	Rock	Metal or Hardrock	Punk	I
1005	5.0	2.0	5.0	2.0	2.0	5.0	4.0	4.0	4.0	3.0	2.0	
1006	4.0	4.0	5.0	1.0	3.0	4.0	1.0	4.0	1.0	1.0	4.0	
1007	4.0	3.0	1.0	1.0	2.0	2.0	2.0	3.0	4.0	1.0	2.0	
1008	5.0	3.0	3.0	3.0	1.0	3.0	1.0	3.0	4.0	1.0	1.0	
1009	5.0	5.0	4.0	3.0	2.0	3.0	3.0	4.0	1.0	1.0	2.0	

```
In [6]: #check shape
data.shape
```

```
Out[6]: (1010, 150)
```

```
In [7]: #check more info mation
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1010 entries, 0 to 1009
Columns: 150 entries, Music to House - block of flats
dtypes: float64(134), int64(5), object(11)
memory usage: 1.2+ MB
```

```
In [8]: data.describe()
```


```
Out[8]:
```

	Music	Slow songs or fast songs	Dance	Folk	Country	Classical music	House - block of flats
count	1007.000000	1008.000000	1006.000000	1005.000000	1005.000000	1003.000000	1008.000000
mean	4.731877	3.328373	3.113320	2.288557	2.123383	2.956132	2.956132
std	0.664049	0.833931	1.170568	1.138916	1.076136	1.252570	1.252570
min	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	5.000000	3.000000	2.000000	1.000000	1.000000	2.000000	2.000000
50%	5.000000	3.000000	3.000000	2.000000	2.000000	3.000000	3.000000
75%	5.000000	4.000000	4.000000	3.000000	3.000000	4.000000	4.000000
max	5.000000	5.000000	5.000000	5.000000	5.000000	5.000000	5.000000

Exploratory analysis

Missing values

Out[9]: <AxesSubplot:>



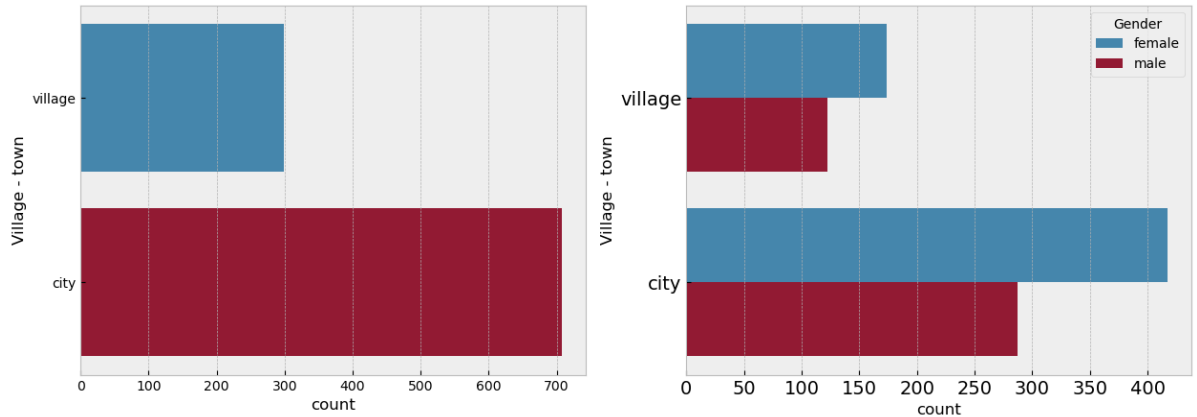
```
Number of people with omitted weight or height: 30
Number of fields that were omitted by people who did not fill Weig
ht or Height: 18
```

Understanding our goal

```
In [12]: var_of_interest = 'Village - town'
mapping = {var_of_interest: {'city': 0, 'village': 1}}
data.dropna(subset=[var_of_interest], inplace=True)
# to be able to use hue parameter for better comparison in seaborn
data["all"] = ""
```

```
In [14]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(15, 5))
sns.countplot(y=var_of_interest, data=data, ax=ax[0])
sns.countplot(y=var_of_interest, hue='Gender', data=data, ax=ax[1])
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
```

```
Out[14]: (array([0, 1]), [Text(0, 0, 'village'), Text(0, 1, 'city')])
```




Outliers

```
In [15]: fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(20,5))
data = data.dropna(subset=['Height'])
sns.violinplot(x='Height', y = "all", hue=var_of_interest, data=data)
data = data.dropna(subset=['Weight'])
sns.violinplot(x='Weight', y = "all", hue=var_of_interest, data=data)

var_of_int_ser = data[var_of_interest]
sns.distplot(data[var_of_int_ser=='village'].Age.dropna(),
              label='village', ax=ax[2], kde=False, bins=30);

sns.distplot(data[var_of_int_ser=='city'].Age.dropna(),
              label='city', ax=ax[2], kde=False, bins=30);
ax[2].legend()
```

```
Out[15]: <matplotlib.legend.Legend at 0x7fe0c7263ca0>
```



As we see there are some outliers that disturb the visualisation.

```
In [16]: display(data[data['Height']<70][['Age', 'Height', 'Weight', 'Gender']])
display(data[data['Weight']>120][['Age', 'Height', 'Weight', 'Gender']])
```

	Age	Height	Weight	Gender	Village - town
676	20.0	62.0	55.0	female	city


	Age	Height	Weight	Gender	Village - town
859	20.0	190.0	125.0	male	city
992	30.0	200.0	150.0	male	city

```
In [18]: data.drop([676, 992, 859], inplace = True)
```

Interestingly, there is a small second hill in Height in city people around 185 cm. The horizontal lines are quartiles.

As TheTS mentioned in the comments, we could look at BMI of rural versus urban people. BMI is calculated as follows: $\text{weight}/\text{height}^2$. The hypothesis is that urban people will have a lower BMI as they might spend more times outdoors.


```
In [20]: data['BMI'] = round(data['Weight']/((data['Height']/100)**2),1)
fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(15,5))
data = data.dropna(subset=['BMI'])
sns.violinplot(x='BMI', hue=var_of_interest, y='Gender', data=data,
               split=True, inner='quartile', ax=ax);
```




```
In [21]: import scipy.stats as stats
city_bmi = data[data[var_of_interest]=='city'].BMI
village_bmi = data[data[var_of_interest]=='village'].BMI
t, p = stats.ttest_ind(village_bmi, city_bmi, axis=0, equal_var=False)
print(' t-stat = {t} \n p-value = {p}'.format(t=t,p=p/2))
```

```
t-stat = 1.7734182239050904
p-value = 0.03837342374443175
```

```
In [22]: fig, ax = plt.subplots(nrows=2, ncols=1, figsize=(15,6))
data_under = data[data['Age']<20]
data_above = data[data['Age']>=20]
sns.violinplot(x='BMI', hue=var_of_interest, y='Gender', data=data_under,
               inner = 'quartile', ax=ax[0], hue_order=['village', 'city'])
sns.violinplot(x='BMI', hue=var_of_interest, y='Gender', data=data_above,
               inner = 'quartile', ax=ax[1], hue_order=['village', 'city'])
```



```
In [23]: fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(15,6))
sns.violinplot(x='BMI', hue=var_of_interest, y='Gender', data=data_under,
               split=True, inner='stick', ax=ax, hue_order=['village', 'city'])
```



Interesting differences

In this section we will analyze differences between the individuals based on the area of living.

Correlation

Firstly, look at correlations between the characteristics and the urban-rural area. Correlation describes the degree of relationship between two variables. However, it tells nothing about the causality. Just a small example, the anti-violent gaming policies say that there is a correlation between time spent on playing violent computer games and a violent behaviour. In fact, we do not know if a the computer games make a person violent ora violent person would play more violent games.

(Btw, I am a big fan of long-lasting code and functions that can be applied on different datasets :)). So, the function `correlation_plot` can be applied on different datasets. The function produces two plots: one for numerical features, another for categorical.


```
In [24]: def do_plotting(x, y, figsize):
fig, ax = plt.subplots(figsize=figsize)
ax.set_title("Correlation coefficient of the variables")
sns.barplot(x=x, y=y, ax=ax)
ax.set_ylabel("Correlation coefficients")

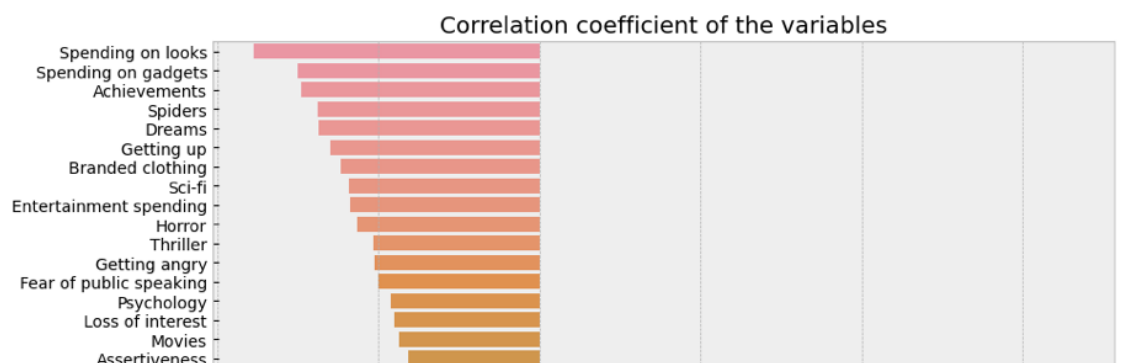
def correlation_plot(var_of_interest, df_main, mapping, figsize=(10
def calc_corr(var_of_interest, df, cols, figsize):
    lbls = []
    vals = []
    for col in cols:
        lbls.append(col)
        vals.append(np.corrcoef(df[col], df[var_of_interest])[0
    corrs = pd.DataFrame({'features': lbls, 'corr_values': vals
    corrs = corrs.sort_values(by='corr_values')
    do_plotting(corrs.corr_values, corrs['features'], figsize)
    return corrs


#imputing the set
df = copy.deepcopy(df_main)
df.replace(mapping, inplace=True)
mean_values = df.mean(axis=0)
df.fillna(mean_values, inplace=True)

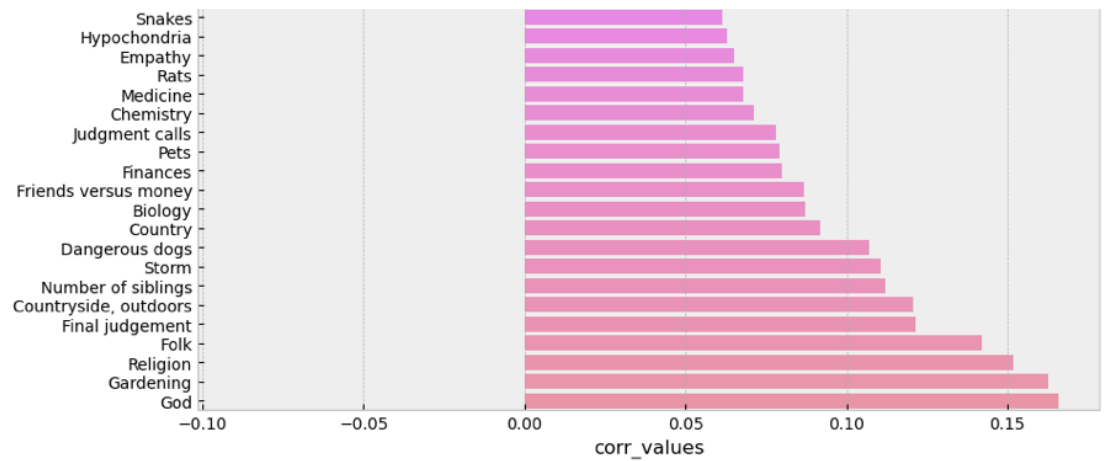
#correlating non-categorical variables
cols_floats = [col for col in df.columns if df[col].dtype != 'o
cols_floats.remove(var_of_interest)
corrs_one = calc_corr(var_of_interest, df, cols_floats, figsize

#correlating categorical variables
cols_cats = [col for col in df.columns if df[col].dtype == 'obj
if cols_cats:
    df_dummies = pd.get_dummies(df[cols_cats])
    cols_cats = df_dummies.columns
    df_dummies[var_of_interest] = df[var_of_interest]
    corrs_two = calc_corr(var_of_interest, df_dummies, cols_cat
else:
    corrs_two = 0
return [corrs_one, corrs_two]
```


```
In [26]: corrs_area = correlation_plot(var_of_interest, data, mapping)
```







Correlation coefficient of the variables



The strongest correlations that we have are coming from the house type and it is quite logical because people in the village would live most of the time in the houses. Other correlations that are not that strong are the associations with God and Spending on looks. We will dig into it.

```
In [27]: #The strongest correlations that we have are
corr_num = corrs_area[0]
corr_cats = corrs_area[1]
display(corr_num[corr_num.corr_values == max(corr_num.corr_values)])
display(corr_num[corr_num.corr_values == min(corr_num.corr_values)])
display(corr_cats[corr_cats.corr_values == max(corr_cats.corr_value)])
display(corr_cats[corr_cats.corr_values == min(corr_cats.corr_value)])
```

	features	corr_values
101	God	0.165819

	features	corr_values
132	Spending on looks	-0.088743

	features	corr_values
31	House - block of flats_house/bungalow	0.600989

	features	corr_values
30	House - block of flats_block of flats	-0.594615

Characteristic differences

I have picked the features that were different among people from urban and rural areas. The plot of all features can be found at the end of this notebook.

In [28]:


```

good_columns = ['Dance', 'Folk', 'Techno, Trance', 'Religion', 'Medicine', 'Spending on gadgets', 'Hypochondria', 'Western', 'Eating to survive', 'God', 'Chemistry', 'Gardening', 'Politics', 'Economy Management', 'Branded clothing', 'Friends versus money', 'Number of siblings', 'Snakes', 'Storm', 'Rats', 'Country', 'Dangerous dogs', 'Finances', 'Entertainment spending', 'Horror', 'Pets', 'Priori', 'Biology', 'Final judgement', 'Sci-fi', 'Spending on gadgets']

fig, ax = plt.subplots(nrows=5, ncols=7, figsize=(30,40), sharex=True)
start = 0
for j in range(5):
    for i in range(7):
        if start == len(good_columns):
            break
        sns.barplot(y=good_columns[start], x=var_of_interest, data=data)
        ax[j,i].set_ylabel('')
        ax[j,i].set_xlabel('')
        ax[j,i].set_title(good_columns[start], fontsize=25)
        start += 1

```





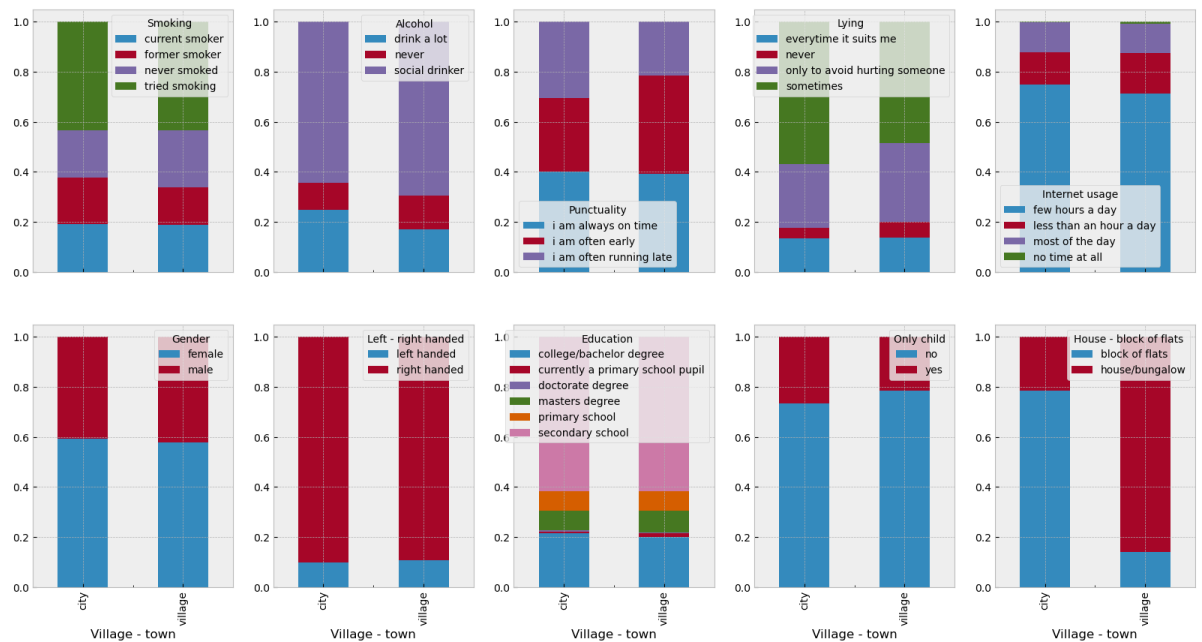
There are several interesting observations:

People from rural area are more into folk music; are more likely to spend time outdoors gardening; also there are more interested in biology, medicine and chemistry. It is quite logical as people in villages are often more close to the nature and therefore would spend time outdoors exploring. Also, the phobias of snakes, rats, dangerous dogs and storm are well-marked in people living in rural areas as they could have had a direct contact to it and know about the circumstances. On the other hand, they are less afraid of the spiders which again can be explained due to the fact that they were more exposed to the nature. Village-livers are more hypochondriac than city-livers which can be explained that there are less hospitals and doctors in villages.

What is a bit striking is that people from rural areas are more religious (Religion and God) than in the city. Additionally, they are prone to "Final judgement" aka "I believe that bad people will suffer one day and good people will be rewarded". It can also have something to do with the stronger belief in God.

City-livers would spend more money on looks, entertainment and branded clothing as normally cities offer much more possibilities to go shopping and have more entertainment facilities than the villages. Also, they would spend more time on gadgets.

```
In [31]: features_cats = [col for col in data.columns if data[col].dtype=='o']
features_cats.remove(var_of_interest)
fig, ax = plt.subplots(nrows=2, ncols=5, figsize=(20,10), sharex=True)
start = 0
for j in range(2):
    for i in range(5):
        tab = pd.crosstab(data[var_of_interest], data[features_cats])
        tab_prop = tab.div(tab.sum(1).astype(float), axis=0)
        tab_prop.plot(kind="bar", stacked=True, ax=ax[j,i])
        start += 1
```



In the categorial features only one thing catches the eye, namely, that village-livers mst of the time live in the houses comparing to the city livers who most live in the flat. That is also self-explanatory.

Multicollinearity

Apart from outliers and missing values, multicollinearity is another common issue. Some ML algorithms like Random Forest do not suffer from multicollinearity, whereas linear regression could have problems with it.

So, it would be quite exciting to look which characteristics are correlated in our data set.

```
In [32]: = data.corr()
e: https://stackoverflow.com/questions/17778394/list-highest-correlation-matrix-is-symmetric-so-we-need-to-extract-upper-triangle-matrix-with

(corr.where(np.triu(np.ones(corr.shape), k=1).astype(np.bool))
         .stack()
         .sort_values(ascending=False))
ay(os.head(10))
ay(os.tail(10))
```

Weight	BMI	0.845273
Height	Weight	0.737569
Biology	Medicine	0.724598
	Chemistry	0.688375
Fantasy/Fairy tales	Animated	0.676642
Shopping	Shopping centres	0.649102
Chemistry	Medicine	0.632300
Classical music	Opera	0.600121
Mathematics	Physics	0.588532
Snakes	Rats	0.565351

dtype: float64

Action	Life struggles	-0.310373
Reading	Cars	-0.317057
Romantic	Height	-0.318037
Cars	Life struggles	-0.323088
Loneliness	Energy levels	-0.341980
Changing the past	Happiness in life	-0.357691
Dangerous dogs	Small - big dogs	-0.376013
Life struggles	Weight	-0.380090
	Height	-0.401332
Loneliness	Happiness in life	-0.440306

dtype: float64


The most correlated features are Height and Weight, what again self-explanatory. People who are interested in Biology are also interested in Medicine and Chemistry. The same for Fantasy/Fairy tales and Animated movies.

You might ask why there is a negative correlation between Life struggles, Weight and Height. I will disclose it at the end ;) and will confront once again with the definition of correlation.

Meanwhile let's have a look on the correlations among our good features.


```
In [33]: corr = data[good_columns].corr()
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(
    corr,
    mask=mask,
    cmap=cmap,
    vmax=.3,
    center=0,
    square=True,
    linewidths=.5,
    cbar_kws={"shrink": .5})
```

Out [33]: <AxesSubplot:>



There is indeed a correlation between Final Judgement and God! Religious people tend to believe more that the bas ones will suffer.

```
In [34]: print (os['Final judgement'][0:2])
```

```
God          0.491327
Charity      0.174467
dtype: float64
```

```
In [35]: print (os['Religion'][0:2])
```

```
God          0.508850
Final judgement 0.289225
dtype: float64
```

The ones who spend money on entertainment, also spend money on looks and gadgets.

```
In [36]: print (os['Entertainment spending'])
```

```
Spending on looks          0.402580
Spending on gadgets        0.336548
Height                    0.160816
Spending on healthy eating 0.143228
Weight                    0.131049
BMI                       0.064862
Number of siblings         0.031045
Age                       -0.042523
dtype: float64
```

Lets exclude of the element from the correlation pair once the correlation is more than 0.5

```
In [37]: corr = data.corr()
os = (corr.where(np.triu(np.ones(corr.shape), k=1).astype(np.bool))
      .stack()
      .sort_values(ascending=False))
display(os[abs(os)>0.5])
drop_colinera_cols = os[abs(os)>0.5].reset_index()['level_1']
```

Weight	BMI	0.845273
Height	Weight	0.737569
Biology	Medicine	0.724598
	Chemistry	0.688375
Fantasy/Fairy tales	Animated	0.676642
Shopping	Shopping centres	0.649102
Chemistry	Medicine	0.632300
Classical music	Opera	0.600121
Mathematics	Physics	0.588532
Snakes	Rats	0.565351
Art exhibitions	Theatre	0.548380
Metal or Hardrock	Punk	0.543569
Rock	Metal or Hardrock	0.526920
Fear of public speaking	Public speaking	0.509547
Religion	God	0.508850
Horror	Thriller	0.508406
Shopping	Spending on looks	0.506264
Rock	Punk	0.504536

dtype: float64

Preparing the dataset for ML

```
In [38]: clean_data = data.dropna(subset=[var_of_interest])
features_int = [col for col in clean_data.columns if clean_data[col]
features_cats = [col for col in clean_data.columns if clean_data[col]

features_int = list(set(features_int) - set(drop_colinera_cols))
print('Number of features {:.0f}'.format(len(features_int)))
```

Number of features 124

We will impute missing values with the mean, although there are some better solutions to do it, like imputing Height and Weight according to the Gender or taking randomly a value in the range [mean - std, mean + std]

```
In [39]: X = clean_data[features_int]
mean_values = X.mean(axis=0)
X = X.apply(lambda x: x.fillna(x.mean()),axis=0)
#X_cats = clean_data[features_cats].drop(var_of_interest, 1)
#X_cats = X_cats.drop('House - block of flats', 1)
#X_cats = pd.get_dummies(X_cats)
#print(X.shape)
#print(X_cats.shape)
```

```
In [41]: from sklearn.metrics import make_scorer, accuracy_score, roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

Y = clean_data[var_of_interest]
for key, val in mapping[var_of_interest].items():
    Y.replace(key,val, inplace = True)
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size
```

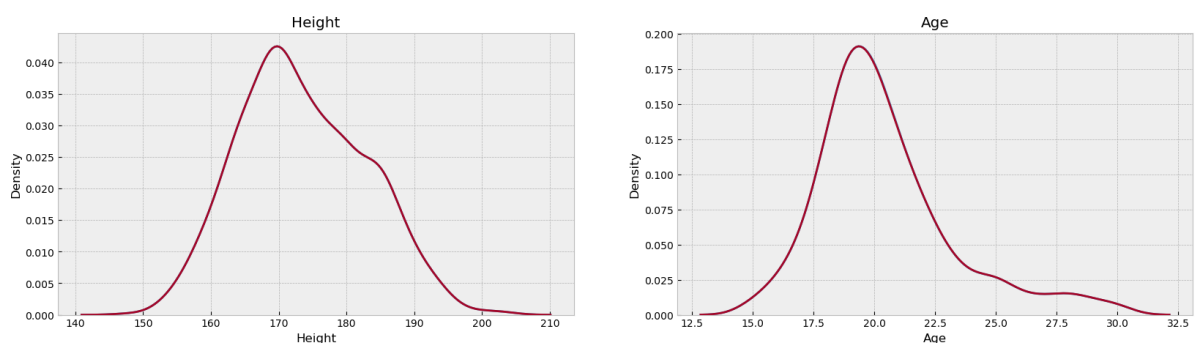
It is also a good idea to have a look on the distribution once we imputed the values to be sure that we did not disrupt it.

```
In [42]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20,5))

sns.kdeplot(X.Height,label = 'Before imputation', ax = ax[0]);
sns.kdeplot(clean_data.Height, label = 'After imputation', ax = ax[0]);
ax[0].set_title('Height');

sns.kdeplot(X.Age,label = 'Before imputation', ax = ax[1]);
sns.kdeplot(clean_data.Age, label = 'After imputation', ax = ax[1])
ax[1].set_title('Age');

#sns.kdeplot(X.Weight,label = 'Before imputation', ax = ax[2])
#sns.kdeplot(clean_data.Weight, label = 'After imputation' , ax = ax[2])
#ax[2].set_title('Weight')
```



Machine Learning

We will standardize variables to be sure that everything is on whie scale as it plays quite an important role in regularization as we will apply logistic regression. The goal is to find variables that have more effect on the dependent variable (aka rural or urban liver).

(One also could use another metric to optimize like f1 but in fact accuracy metrix performed somehow better on this dataset regarding the f1 and accuracy score on test set.)

```
In [43]: # standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)
```

```

In [47]: # gridsearch for parameter tuning
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.linear_model import LogisticRegression
clr = LogisticRegression()
KF = KFold(len(x_train))
param_grid = {'C': [.001, .01, .03, .1, .3, 1, 3, 10]}
grsearch = GridSearchCV(clr, param_grid=param_grid, cv=KF, scoring
grsearch.fit(x_train, y_train)
print(grsearch.best_params_)

# fitting logistic regression and evaluating
clr = LogisticRegression(C=grsearch.best_params_['C'])
clr.fit(x_train, y_train)

mean_accuracy = np.mean(cross_val_score(clr, x_train, y_train, cv=K
print('Average accuracy score on CV set: {:.2f}'.format(mean_accura

mean_f1 = np.mean(cross_val_score(clr, x_train, y_train, cv=KF, sco
print('Average f1 on CV set: {:.2f}'.format(mean_f1))
print('')
print('Accuracy score on test set is: {:.2f}'.format(clr.score(x_te
recall = recall_score(y_test, clr.predict(x_test))
print('Recall on test: {:.2f}'.format(recall))
precision = precision_score(y_test, clr.predict(x_test))
print('Presicion on test: {:.2f}'.format(precision))
print('F1 score on test: {:.2f}'.format((2*recall*precision)/(reca

```

```
{'C': 3}
```

```
Average accuracy score on CV set: 0.67
```

```
Average f1 on CV set: 0.08
```

```
Accuracy score on test set is: 0.70
```

```
Recall on test: 0.40
```

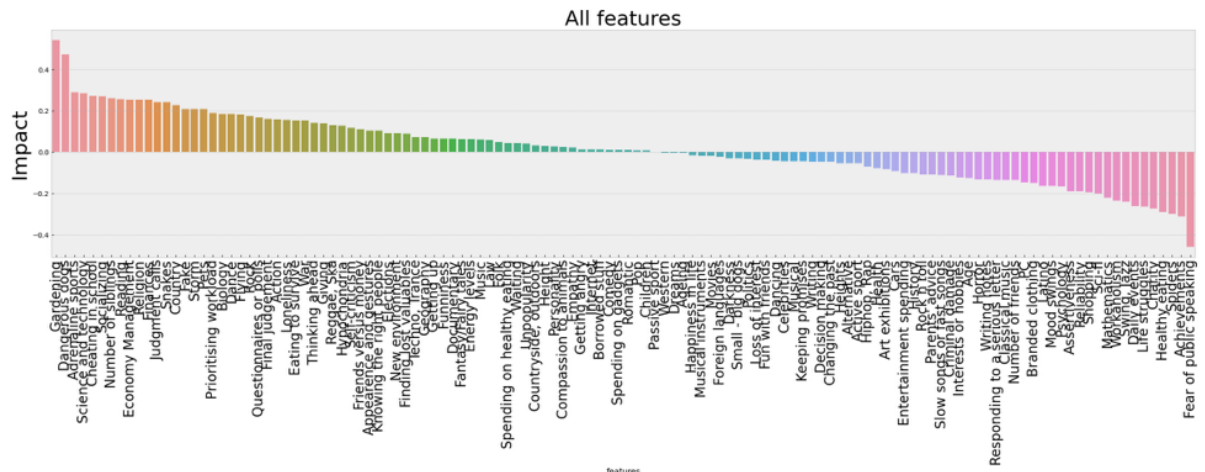
```
Presicion on test: 0.51
```

```
F1 score on test: 0.45
```

Lets look at the impact of all features on our rural-village liver classification. And no worries, we will zoom in.

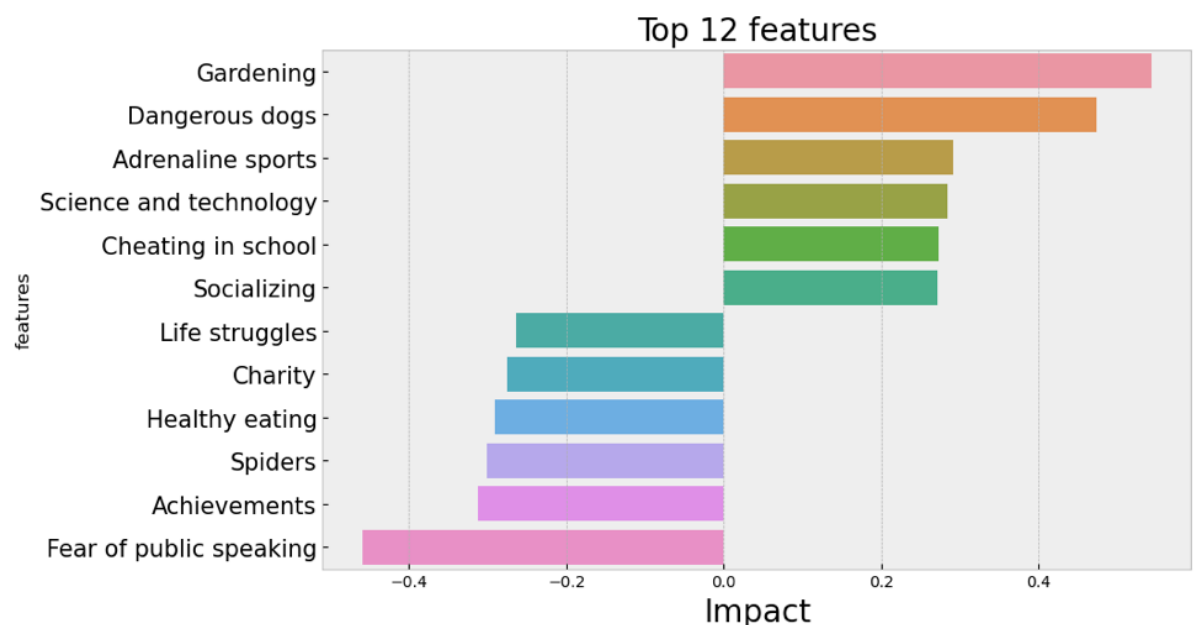
```
In [48]: feat_coeff = pd.DataFrame({'features': X.columns, 'impacts': clr.coef})
feat_coeff = feat_coeff.sort_values('impacts', ascending=False)

fig, ax1 = plt.subplots(1,1, figsize=(30,6));
sns.barplot(x=feat_coeff.features, y=feat_coeff.impacts, ax=ax1);
ax1.set_title('All features', size=30);
ax1.set_xticklabels(labels=feat_coeff.features, size=20, rotation=90);
ax1.set_ylabel('Impact', size=30);
```



Zooming on only top 12 features based on their coefficients:

```
In [49]: top10 = pd.concat([feat_coeff.head(6), feat_coeff.tail(6)])
fig, ax1 = plt.subplots(1,1, figsize=(10,6))
sns.barplot(y=top10.features, x=top10.impacts, ax=ax1);
ax1.set_title('Top 12 features', size=20);
ax1.set_yticklabels(labels=top10.features, size=15);
ax1.set_xlabel('Impact', size=20);
```



You can interpret it as : For every unit change in gardening (so, you thought ok, it is actually 5 not 4) the log odds of you having lived in city is decreased by 0.156.

It will get a bit simpler if we exponentiate the coefficients. The interpretation would be then: For every unit change in gardening the odds of you having lived in city is decreased by a factor $\exp(-0.156) = 0.8555592$ (or by 85%).

Disclaimer: As we have quite lots of features and only 1K samples, it would be good to do somekind of feature selection as linear models tend to suffer from non-informative features. (see Max Kuhn "Applied Predictive Modeling", Feature Selection)

Insights:

- Gardening, fear of dogs and storm, as well as religion decrease the odds of having lived in the city. Rural area livers are more exposed to gardening and might have more experience with dangerous dogs. As well as storm might have a more devastating effect on the village than on a city. People in the village are also typically a bit more religious than in the city.
- The fear of spiders increase the odds of living in the city. It is also intuitive as city livers are less used to spiders than the village livers.

Sci-fi genre of movies might, healthy eating and fear of public speaking as well as getting angry are a bit speculative. Do urban livers get angry faster also why do they prefer Sci-fi more?

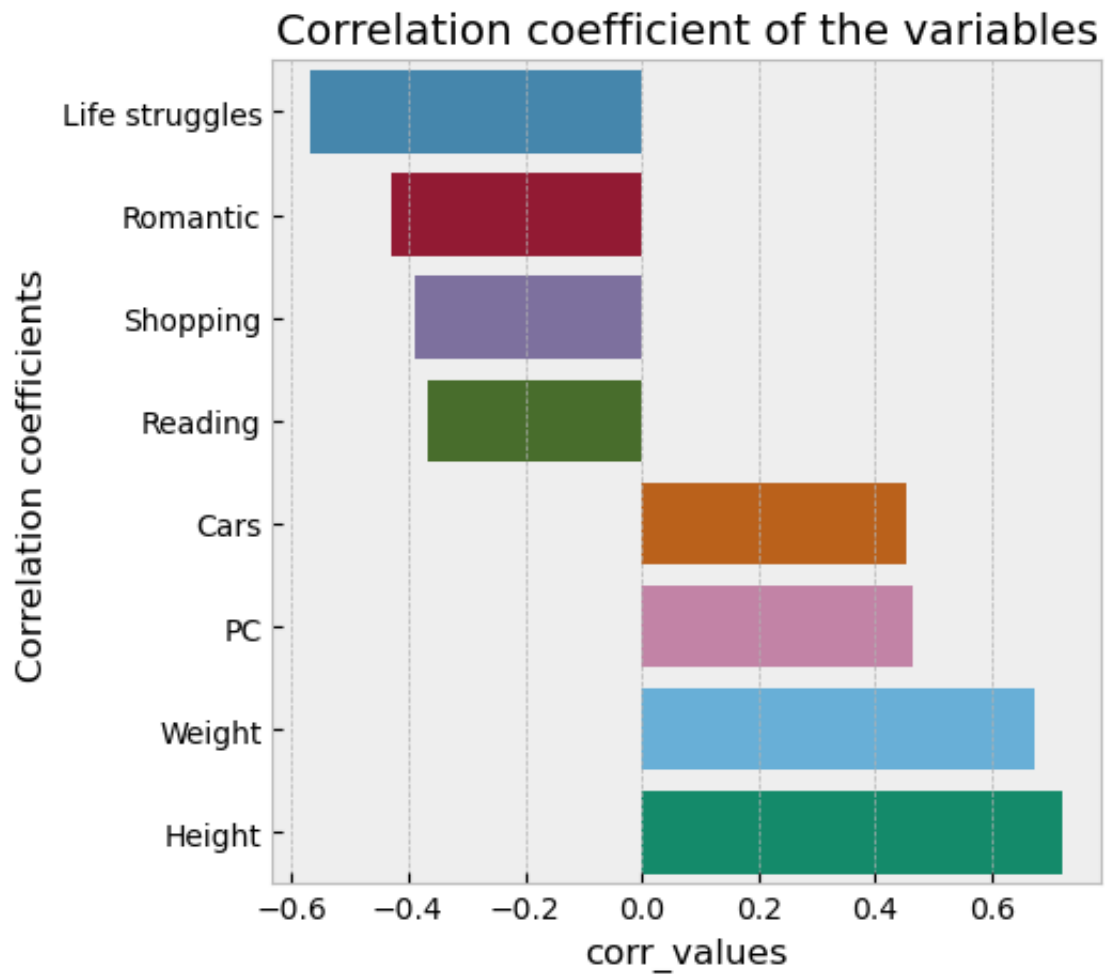
The dessert: negative correlation between Life struggles, Weight and Heigt

```
In [50]: display(os.tail(3))
```

```
Life struggles  Weight          -0.380090
                Height          -0.401332
Loneliness     Happiness in life -0.440306
dtype: float64
```

To understand why there is a negative correlation between life struggles, weight and height I will change the objectives. Now, we want to analyze the data based on the gender difference. So, lets plot correlation of each variable against the gender.


```
In [52]: cols_to_keep = ['Life struggles', 'Romantic', 'Shopping',  
                        'Reading', 'Weight', 'Height', 'PC',  
                        'Cars', 'Gender']  
gender_map = {'Gender': {'female': 0, 'male': 1}}  
corrs_dfs_gender = correlation_plot('Gender', data[cols_to_keep], g
```




And a regression that will help us to get more insights.

```
In [53]: clean_data = data.dropna(subset=['Gender'])
features_int = [col for col in clean_data.columns if clean_data[col]
X = clean_data[features_int]
mean_values = X.mean(axis=0)
X = X.apply(lambda x: x.fillna(x.mean()),axis=0)
Y = clean_data['Gender']
Y.replace('female',0, inplace = True)
Y.replace('male',1, inplace = True)

scaler = StandardScaler()
scaler.fit(X)
X = scaler.transform(X)

clr = LogisticRegression()
clr.fit(X, Y)
feat_coeff = pd.DataFrame({'features': features_int, 'impacts': clr.
feat_coeff = feat_coeff.sort_values('impacts', ascending=False)

top10 = pd.concat([feat_coeff.head(8), feat_coeff.tail(8)])
fig, ax1 = plt.subplots(1,1, figsize=(10,6))
sns.barplot(y=top10.features, x=top10.impacts, ax=ax1);
ax1.set_title('Top 16 features', size=20);
ax1.set_yticklabels(labels=top10.features, size=15);
ax1.set_xlabel('Impact', size=20);
```



Additional plots

In [55]:


```


fig, ax = plt.subplots(nrows=15, ncols=8, figsize=(30, 70), sharex=
start = 0
for j in range(15):
    for i in range(8):
        sns.barplot(
            y=features_int[start], x=var_of_interest, data=data, ax
            ax[j, i].set_ylabel('')
            ax[j, i].set_xlabel('')
            ax[j, i].set_title(features_int[start], fontsize=25)
            start += 1

fig, ax = plt.subplots(nrows=1, ncols=4, figsize=(7, 2), sharex=True

for i in range(4):
    sns.barplot(y=features_int[start], x=var_of_interest, data=data
    ax[i].set_ylabel('')
    ax[i].set_xlabel('')
    ax[i].set_title(features_int[start], fontsize=10)
    start += 1

```

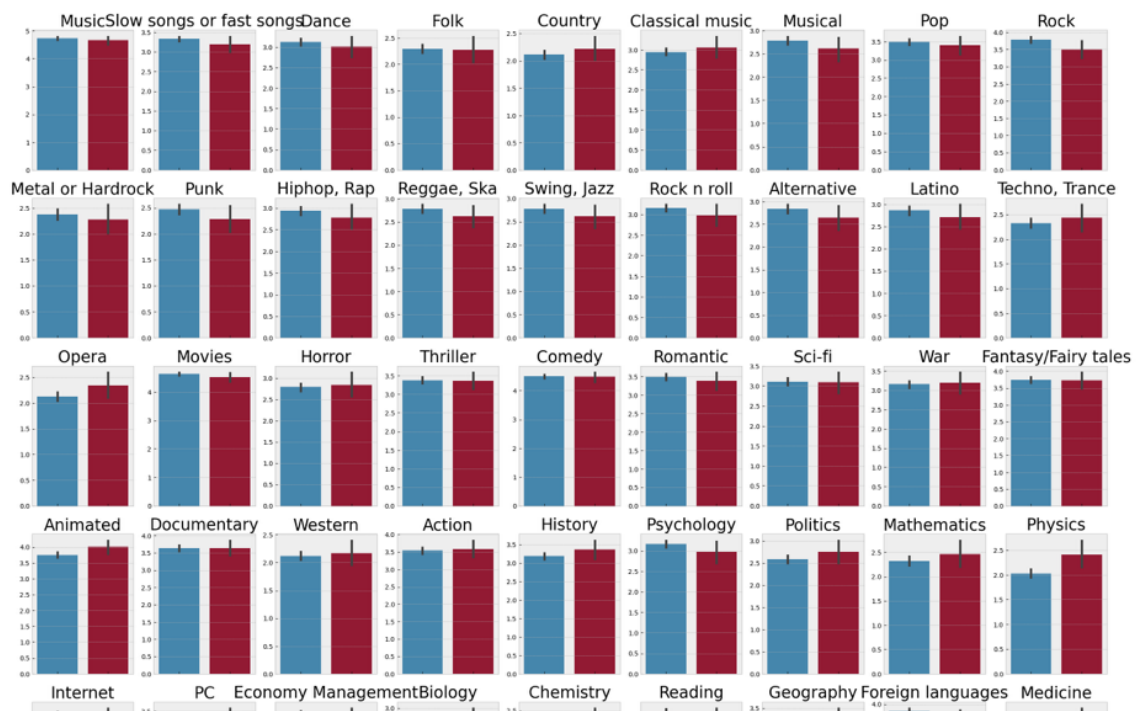




Initially I was intended to analyze the difference between left and high handed but the class distribution is even more imbalanced.

```
In [56]: fig, ax = plt.subplots(nrows=15, ncols=9, figsize=(30, 70), sharex=
start = 0
for j in range(15):
    for i in range(9):
        sns.barplot(
            y=features_int[start],
            x='Left - right handed',
            data=data,
            ax=ax[j, i])
        ax[j, i].set_ylabel('')
        ax[j, i].set_xlabel('')
        ax[j, i].set_title(features_int[start], fontsize=25)
        start += 1

fig, ax = plt.subplots(nrows=1, ncols=4, figsize=(7, 2), sharex=True)
for i in range(4):
    sns.barplot(
        y=features_int[start], x='Left - right handed', data=data,
        ax[i].set_ylabel('')
        ax[i].set_xlabel('')
        ax[i].set_title(features_int[start], fontsize=10)
        start += 1
```



In []: